

Rotman

INTRO TO SQL

BTA SQL Workshop (<https://tdmdal.github.io/sql-bta-2021/>)

October 18, 2021 Prepared by Jay / TDMDAL



Rotman School of Management
UNIVERSITY OF TORONTO

Goal for Today (2 hrs)


- Understand what's SQL and related concepts such as
 - Database (DB) and relational DB (RDB)
 - DB management system (DBMS) and RDBMS
- Get a taste of SQL coding
 - Simple column and row operations
 - Simple aggregation
 - simple join operation

What's SQL (Structured Query Language)

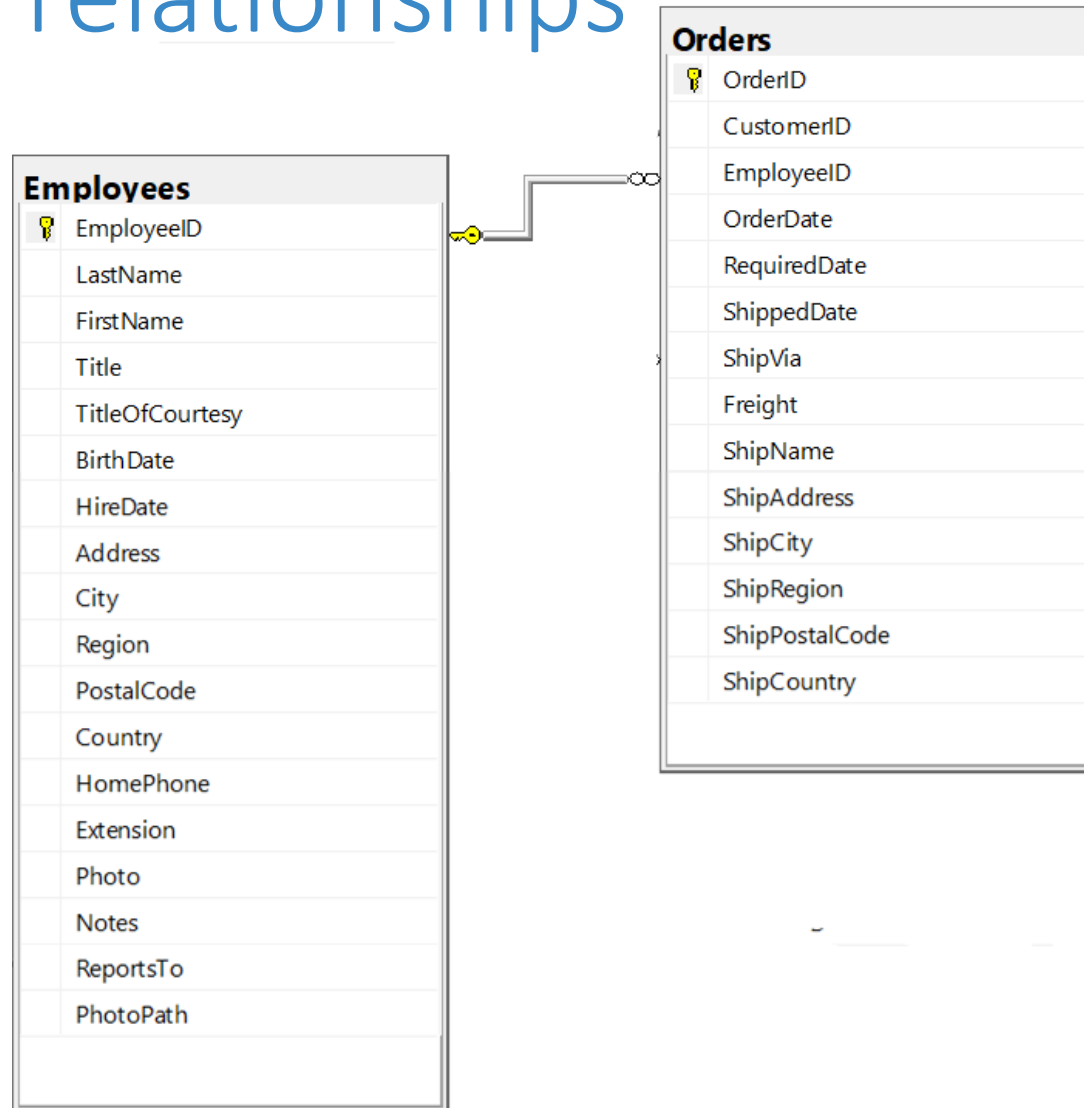
- Most widely used database (DB) language
 - a domain specific language (managing data stored in relational DBs)
- Not a proprietary language
 - Open specifications/standards
 - All major DBMS (DB Mgmt. System) vendors implement ANSI Standard SQL
 - However, SQL Extensions are usually DB specific (SQL dialects)
- Powerful despite simplicity

What's a DB and a Relational DB

- What's a database: A collection of data in an organized way
- *Relational DB (RDB)*
 - tables
 - columns/fields/variables and datatypes
 - rows/records/observations
 - primary key, foreign key, constraints and relationships
 - other objects: indices, views, and many more




Primary key, foreign key, constraints and relationships



EmployeeID	LastName	FirstName	Title	...
1	Davolio	Nancy	Sales Representative	...
2	Fuller	Andrew	Vice President, Sales	...
3	Leverling	Janet	Sales Representative	...
4	Peacock	Margaret	Sales Representative	...
...

OrderID	CustomerID	EmployeeID	...
10248	VINET	5	...
10249	TOMSP	6	...
10250	HANAR	4	...
...

Northwind DB

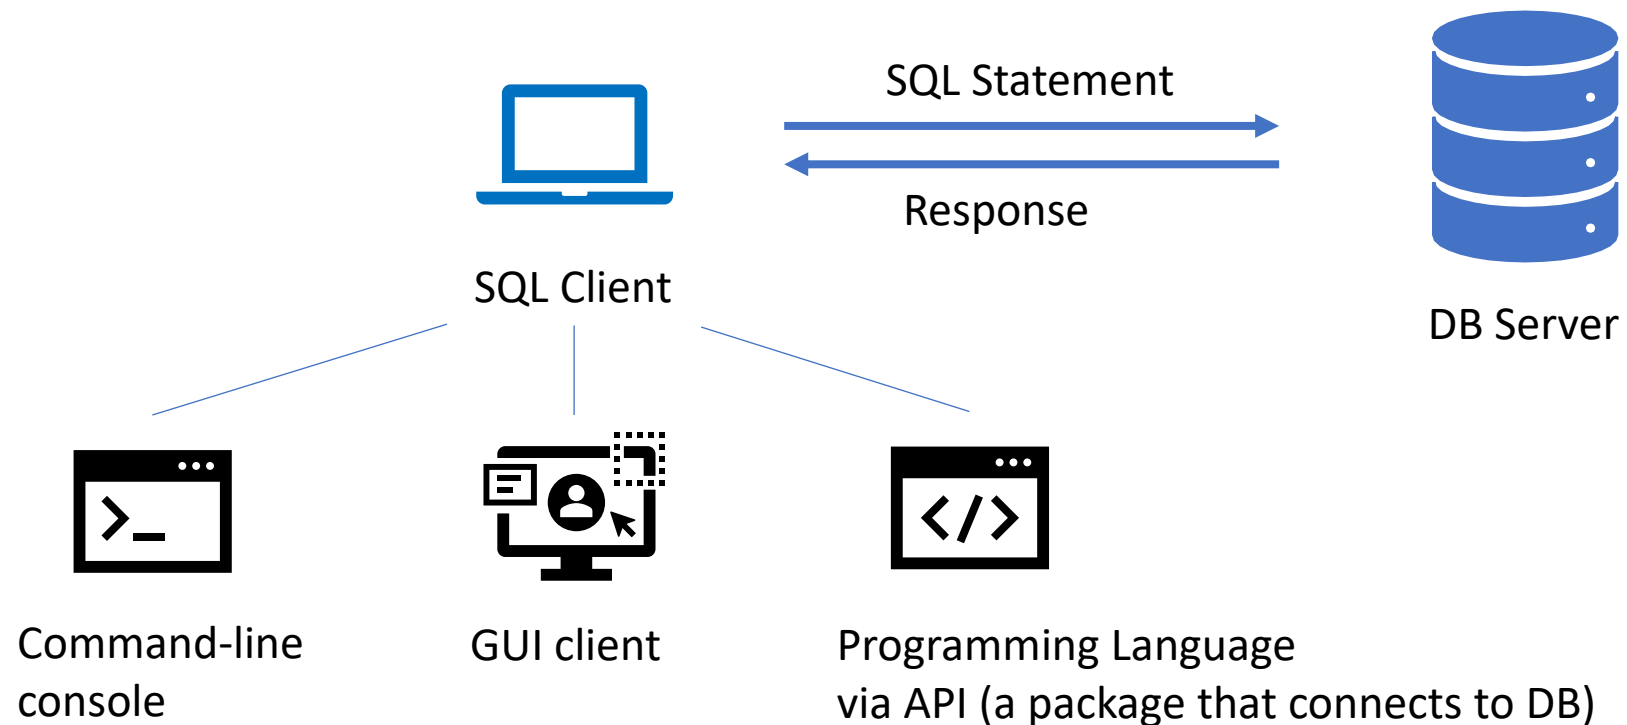


What is a DB Management System

- A software system that manages/maintains DBs
- A few examples of Relational DBMS (RDBMS)
 - Open source: SQLite, MariaDB, PostgreSQL
 - Commercial: MySQL, Microsoft SQL Server, Oracle, etc.






Connect to a DB and write SQL – Architecture



Note: SQL client and DB server can be on the same computer

Connect to a DB and use SQL – SQL Clients

- DB specific management client
 - command-line console
 - GUI (Graphic User Interface) client
 - e.g., [DB Browser for SQLite](#), [MySQL Workbench](#), [pgAdmin for PostgreSQL](#), [MS SSMS](#)
- Generic DB client can connect to different DBs through connectors
 - GUI client (e.g. [DBeaver](#), [Beekeeper Studio](#), [Navicat](#))   
 - Programming language
 - e.g., Python + [SQLAlchemy](#) + DBAPI (e.g. [SQLite](#), [MySQL](#), [PostgreSQL](#), etc.), R + [dbplyr](#)
 - **In this workshop:** Python + [ipython-sql notebook magic](#) (depends on [SQLAlchemy](#))

Beyond a relational DB language

- SAS's PROC SQL
- Spark's SparkSQL
 - [Apache Spark](#) is a big data computing framework
- Hive's HiveQL, an SQL-like query language
 - [Apache Hive](#) is a distributed data warehouse (data warehouse?)
- **Google BigQuery's SQL** (a great first step to big data analysis)
 - [BigQuery](#) is Google's data warehouse (analyze petabytes of data at ease)

Big Data ML with SQL (e.g. Google BigQuery)

The screenshot displays the Google Cloud Platform BigQuery interface. The top navigation bar includes the Google Cloud Platform logo, a user profile icon, and a search bar. The main interface is divided into three sections: Explorer, Table schema, and SQL editor.

Explorer: Shows a list of pinned projects. The 'new_york_taxi_trips' project is selected, showing a list of tables including 'taxi_zone_geom', 'tlc_fhv_trips_2015', and 'tlc_fhv_trips_2016'.

Table schema: Displays the schema for the 'tlc_fhv_trips_2015' table. The schema includes the following fields:

Field name	Type	Mode	Policy Tags	Description
location_id	INTEGER	NULLABLE		The TLC taxi location ID
pickup_datetime	DATETIME	NULLABLE		The date and time of the pickup
dispatching_base_num	STRING	NULLABLE		The TLC Boro Dispatching Base Number
borough	STRING	NULLABLE		This is the Boro of the pickup location
zone	STRING	NULLABLE		This is the Zone of the pickup location
service_zone	STRING	NULLABLE		This is the Service Zone of the pickup location

SQL editor: Contains the following SQL query:

```
CREATE OR REPLACE MODEL taxi.taxifare_model_2
OPTIONS
  (model_type='linear_reg', labels=['total_fare']) AS
WITH params AS (
  SELECT
    1 AS TRAIN,
    2 AS EVAL
  ),
daynames AS
  (SELECT ['Sun', 'Mon', 'Tues', 'Wed', 'Thurs', 'Fri', 'Sat'] AS daysofweek),
taxitrips AS (
  SELECT
    (tolls_amount + fare_amount) AS total_fare,
    daysofweek[ORDINAL(EXTRACT(DAYOFWEEK FROM
      pickup_datetime))] AS dayofweek,
    EXTRACT(HOUR FROM pickup_datetime) AS hourofday,
    SQRT(POW((pickup_longitude - dropoff_longitude),2) +
      POW((pickup_latitude - dropoff_latitude), 2)) as dist,
    #Euclidean distance between pickup and drop off
    SQRT(POW((pickup_longitude - dropoff_longitude),2))
    as longitude, #Euclidean distance between pickup and drop
    off in longitude
    SQRT(POW((pickup_latitude - dropoff_latitude), 2)) as
```

Ref. [Using BigQuery ML and BigQuery GIS together to predict NYC taxi trip cost](#)

SQL Hands-on Learning (Learning-by-doing)

- Course website: <https://tdmdal.github.io/sql-bta-2021/>
- Google Colab
 - Google's Jupyter Notebook
 - A notebook can contain live code, equations, visualizations and narrative text
- Why SQLite?
 - a [small](#), [fast](#), [self-contained](#), [high-reliability](#), [full-featured](#), SQL DB engine
 - perfect for learning SQL

Hands-on Part 1: Basics

- Retrieve data: `SELECT . . . FROM . . .`
- Sort retrieved data: `SELECT . . . FROM . . . ORDER BY . . .`
- Filter data: `SELECT . . . FROM . . . WHERE . . .`
 - `IN`, `NOT`, `LIKE` and `%` wildcard
- Create calculated fields
 - mathematical calculations (e.g. `+`, `-`, `*`, `/`)
 - data manipulation functions (e.g. `DATE()`, `||`)


Hands-on Part 2: Summarize and Group Data

- Summarize data using aggregate functions (e.g. `COUNT()`, `MIN()`, `MAX()`, and `AVG()`).
- Group data and filter groups: `SELECT . . . FROM . . . GROUP BY . . . HAVING . . .`
- `SELECT` statement syntax ordering
 - `SELECT . . . FROM . . . WHERE . . . GROUP BY . . . HAVING . . . ORDER BY . . .`

Hands-on Part 3: Join Tables

- Inner join: `SELECT...FROM...INNER JOIN...ON...`
- Left join: `SELECT...FROM...LEFT JOIN...ON...`
- Other join variations (see appendix)

Join – Inner Join



```
SELECT *  
FROM Table1  
    INNER JOIN Table2  
    ON Table1.pk = Table2.fk;
```

Table1


pk	t1c1
1	a
2	b

Table2

fk	t2c1
1	c
1	d
3	e

pk	t1c1	fk	t2c1
1	a	1	c
1	a	1	d

Join – Left (Outer) Join



```
SELECT *  
FROM Table1  
  LEFT JOIN Table2  
    ON Table1.pk = Table2.fk;
```

Table1


pk	t1c1
1	a
2	b

Table2

fk	t2c1
1	c
1	d
3	e

pk	t1c1	fk	t2c1
1	a	1	c
1	a	1	d
2	b	null	null

Join - Left (Outer) Join With Exclusion



```
SELECT *  
FROM Table1  
  LEFT JOIN Table2  
    ON Table1.pk = Table2.fk  
WHERE Table2.fk is NULL;
```

Table1

pk	t1c1
1	a
2	b


Table2

fk	t2c1
1	c
1	d
3	e

pk	t1c1	fk	t2c1
2	b	null	null

Learning Resources

- A more complete notebook for SQL exercises on workshop website
- Online resources
 - SQLite Tutorial: <https://www.sqlitetutorial.net/>
 - SQL Tutorial Org: <https://www.sqltutorial.org/>
 - W3 School SQL tutorial: <https://www.w3schools.com/sql/default.asp>
 - [SQL for Data Analysis](#) at Udacity
 - [Learning SQL Programming by Scott Simpson \(1h 27m\)](#) on LinkedIn Learning
- A little book
 - SQL in 10mins a Day (5th edition) by Ben Forta



Appendix

- Many join operation variations
- SQL is much more...

Join – Right Outer Join*




Table1

pk	t1c1
1	a
2	b

Table2

fk	t2c1
1	c
1	d
3	e


```
SELECT *  
FROM Table2  
  LEFT JOIN Table1  
    ON Table2.fk = Table1.pk
```

```
-----  
SELECT *  
FROM Table1  
  RIGHT JOIN Table2  
    ON Table1.pk = Table2.fk;
```

} SQLite doesn't support this RIGHT JOIN key word, but some DBMSs do (e.g. MySQL).

pk	t1c1	fk	t2c1
1	a	1	c
1	a	1	d
null	null	3	e

Join - Right Outer Join With Exclusion*



```
SELECT *  
FROM Table2  
  LEFT JOIN Table1  
    ON Table2.fk = Table1.pk  
WHERE Table1.pk is NULL;
```

```
-----  
SELECT *  
FROM Table1  
  RIGHT JOIN Table2  
    ON Table1.pk = Table2.fk  
WHERE Table1.pk is NULL;
```

Table1

pk	t1c1
1	a
2	b


Table2

fk	t2c1
1	c
1	d
3	e

pk	t1c1	fk	t2c1
null	null	3	e

} SQLite doesn't support this RIGHT JOIN key word, but some DBMSs do (e.g. MySQL).

Join – Full Outer Join



```
SELECT pk, t1c1, fk, t2c1  
FROM Table1  
  LEFT JOIN Table2  
    ON Table1.pk = Table2.fk
```

UNION

```
SELECT pk, t1c1, fk, t2c1  
FROM Table2  
  LEFT JOIN Table1  
    ON Table2.fk = Table1.pk;
```





Table1

pk	t1c1
1	a
2	b

Table2

fk	t2c1
1	c
1	d
3	e

pk	t1c1	fk	t2c1
1	a	1	c
1	a	1	d
2	b	null	null
null	null	3	e

Note: Some DBMS support FULL OUTER JOIN keyword (e.g. MS SQL) so you don't need to do it the above way.

Join – Full Outer Join With Exclusion*








Table1

pk	t1c1
1	a
2	b

Table2

fk	t2c1
1	c
1	d
3	e

```
SELECT pk, t1c1, fk, t2c1
FROM Table1
  LEFT JOIN Table2
    ON Table1.pk = Table2.fk
WHERE Table2.fk is NULL
UNION
SELECT pk, t1c1, fk, t2c1
FROM Table2
  LEFT JOIN Table1
    ON Table2.fk = Table1.pk
WHERE Table1.pk is NULL;
```



pk	t1c1	fk	t2c1
2	b	null	null
null	null	3	e

SQL is much more - 1

- Sub-query
- CTE and temporary table
- Self-join
- CASE keyword
- UNION keyword

SQL is much more - 2

- Insert data (`INSERT INTO...VALUES...; INSERT INTO...SELECT...FROM...`)
- Update data (`UPDATE...SET...WHERE...`)
- Delete data (`DELETE FROM...WHERE...`)
- Manipulate tables (`CREATE TABLE...; ALTER TABLE...; DROP TABLE...`)
- Views (`CREATE VIEW...AS...`)

The list goes on and on

- Stored procedures
- Functions
- Transaction processing
- Cursors (going through table row by row)
- WINDOW function
- Query optimization
- DB permissions & security
- ...

Ref. A stack overflow discussion on [What is “Advanced” SQL.](#)